

Enhancement 5: Choose a new dataset from the list below. Search the Internet and download

- ✓ your chosen dataset (many of them could be available on kaggle). Adapt your model to your dataset. Train your model and record your results.

- cancer\_dataset - Breast cancer dataset.
- crab\_dataset - Crab gender dataset.
- glass\_dataset - Glass chemical dataset.
- iris\_dataset - Iris flower dataset.
- ovarian\_dataset - Ovarian cancer dataset.
- thyroid\_dataset - Thyroid function dataset.

- ✓ Summary of Model Performance:

- **Test Loss:** 0.0059 — The test loss is very low, indicating that the model's predictions are extremely close to the actual values on the test set.
- **Accuracy:** 1.0000 (100%) — The model correctly predicted the labels for all test samples, meaning it made zero classification errors.
- **Precision:** 1.0000 (100%) — Of all the positive predictions made by the model, 100% were correct. There were no false positives.
- **Recall:** 1.0000 (100%) — The model identified all actual positive instances correctly. There were no false negatives.
- **F1-Score:** 1.0000 (100%) — Since both precision and recall are perfect, the F1-score, which balances both, is also 100%.

Conclusion:

- The model performed extremely well on the Iris dataset, achieving **perfect test accuracy** and very low loss values.
- There is no sign of overfitting, as the model performs equally well on both the training and test sets.

```
[29] from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import torch
import torch.nn as nn
import torch.optim as optim

# Load the Iris dataset
iris = load_iris()
X = iris.data # Features
y = iris.target # Labels

# Split into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Scale the features (standardization)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Convert to PyTorch tensors
X_train_tensor = torch.FloatTensor(X_train)
y_train_tensor = torch.LongTensor(y_train)
X_test_tensor = torch.FloatTensor(X_test)
y_test_tensor = torch.LongTensor(y_test)

# Create TensorDatasets and DataLoaders
train_dataset = torch.utils.data.TensorDataset(X_train_tensor, y_train_tensor)
test_dataset = torch.utils.data.TensorDataset(X_test_tensor, y_test_tensor)

train_loader = torch.utils.data.DataLoader(train_dataset, batch_size=32, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=32)
```

```
[26] class IrisModel(nn.Module):
    def __init__(self):
        super(IrisModel, self).__init__()
        # Input size is 4 (features of Iris dataset), output size is 3 (classes)
        self.fc1 = nn.Linear(4, 128) # 4 input features
        self.fc2 = nn.Linear(128, 64)
        self.fc3 = nn.Linear(64, 32)
        self.fc4 = nn.Linear(32, 3) # 3 output classes

    def forward(self, x):
        x = torch.relu(self.fc1(x))
        x = torch.relu(self.fc2(x))
        x = torch.relu(self.fc3(x))
        x = self.fc4(x)
        return x
```

```
[27] # Initialize model, optimizer, and loss function
model = IrisModel()
optimizer = optim.Adam(model.parameters(), lr=1e-3)
criterion = nn.CrossEntropyLoss()

# Training loop
num_epochs = 100

for epoch in range(num_epochs):
    model.train()
    epoch_loss = 0
    epoch_acc = 0
    for X_batch, y_batch in train_loader:
        optimizer.zero_grad()
        outputs = model(X_batch)
        loss = criterion(outputs, y_batch)
        loss.backward()
        optimizer.step()

        epoch_loss += loss.item()

    # Calculate accuracy
    _, preds = torch.max(outputs, 1)
    acc = torch.sum(preds == y_batch).item() / len(y_batch)
    epoch_acc += acc

    print(f'Epoch {epoch+1}/{num_epochs} | Loss: {epoch_loss/len(train_loader):.4f} | Accuracy: {epoch_acc/len(train_loader):.4f}')
```

```
# Evaluate on test data
model.eval()
test_loss = 0
test_acc = 0
with torch.no_grad():
    for X_batch, y_batch in test_loader:
        outputs = model(X_batch)
        loss = criterion(outputs, y_batch)
        test_loss += loss.item()

    # Calculate accuracy
    _, preds = torch.max(outputs, 1)
    acc = torch.sum(preds == y_batch).item() / len(y_batch)
    test_acc += acc

    print(f'Test Loss: {test_loss/len(test_loader):.4f} | Test Accuracy: {test_acc/len(test_loader):.4f}')
```

```
[31] from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score

# Evaluate on the test data
model.eval() # Set the model to evaluation mode
test_loss = 0
all_preds = []
all_labels = []

with torch.no_grad(): # Disable gradient calculation for faster evaluation
    for X_batch, y_batch in test_loader:
        outputs = model(X_batch)
        loss = criterion(outputs, y_batch)
        test_loss += loss.item()

        # Get the predicted class labels
        _, preds = torch.max(outputs, 1)

        # Store the predictions and true labels for metric calculations
        all_preds.extend(preds.cpu().numpy())
        all_labels.extend(y_batch.cpu().numpy())

# Compute accuracy, precision, recall, and F1-score using the collected predictions and true labels
accuracy = accuracy_score(all_labels, all_preds)
precision = precision_score(all_labels, all_preds, average='weighted')
recall = recall_score(all_labels, all_preds, average='weighted')
f1 = f1_score(all_labels, all_preds, average='weighted')

print(f'Test Loss: {test_loss/len(test_loader):.4f}')
print(f'Accuracy: {accuracy:.4f}')
print(f'Precision: {precision:.4f}')
print(f'Recall: {recall:.4f}')
print(f'F1-Score: {f1:.4f}')
```



```
Test Loss: 0.0059
Accuracy: 1.0000
Precision: 1.0000
Recall: 1.0000
F1-Score: 1.0000
```